



Figure 1: This figure shows the selection frequency of each sample at a high pruning ratio (90%). The x-axis represents data indices sorted in descending order by selection frequency, and the y-axis indicates the normalized selection ratio—i.e., how often each data point was selected, divided by the total number of data points seen during training (values sum to 1). Methods like Random Dynamic Pruning, InfoBatch, and IES use nearly the entire dataset during training, making them incomparable to static pruning methods, which are restricted to only 10% of the data.