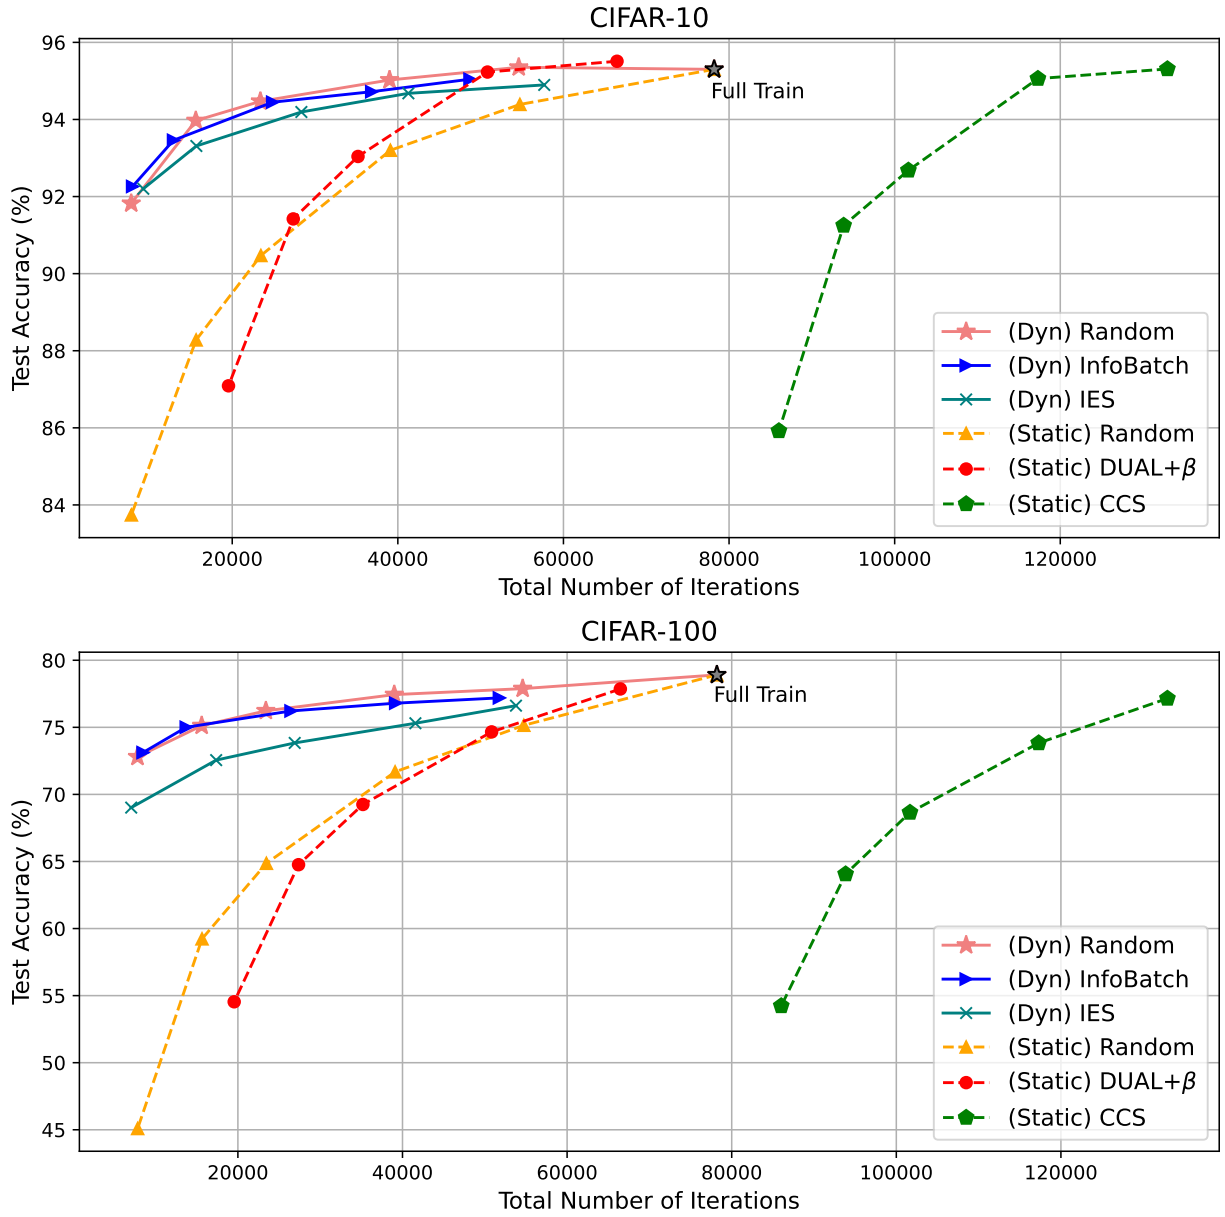

Figure 1: This figure shows test accuracy (y-axis) versus the total number of iterations needed to fully train each subset (x-axis). The top figure corresponds to CIFAR-10, and the bottom to CIFAR-100. Results are averaged over five random seeds. For the static pruning methods, we plotted the number of iterations and test accuracy at pruning ratios of 30%, 50%, 70%, 80%, and 90%, which correspond to the markers on each line from right to left. We adjust the total training epochs of InfoBatch following the procedure in the original paper, in order to evaluate its performance when the total number of iterations is reduced. For IES, we adjust the pruning threshold to evaluate its performance under a reduced total number of training iterations. Results demonstrate that dynamic pruning methods—including even random dynamic pruning—outperform static baselines in accuracy while requiring fewer iterations. Among static methods, DUAL is both the most efficient and the best-performing. However, it does not achieve as favorable a time-performance trade-off as dynamic methods. Still, at low pruning ratios—where more information is available—DUAL performs comparably to dynamic approaches.