



Figure 1: This figure shows test accuracy (y-axis) versus the total number of iterations needed to fully train each subset (x-axis). The top figure corresponds to CIFAR-10, and the bottom to CIFAR-100. Results are averaged over five random seeds. Dynamic pruning methods—including even random dynamic pruning—outperform static baselines in accuracy while requiring fewer iterations. Among static methods, DUAL is both the most efficient and the best-performing. However, it doesn’t achieve as favorable a time-performance trade-off as dynamic methods. Still, at low pruning ratios—where more information is available—DUAL performs comparably to dynamic approaches.